



## **Mastering Rust for Embedded Firmware Development**

**ESE Konference Dec. 2024**





## Embedded Rust Tool Chain

Rust for firmware development

## Aims and Topics

---

In this module we will learn about the various tools used in a Rust based embedded software development environment.

- Minimal environment requirements
- Interactive debugging
- Basic logging



---

# Minimal Environment Requirements



## General Environment Setup with `rustup`

- `rustup` provides the support for:
  - `rustc` / `cargo`
- Install `rustup`
  - <https://rust-lang.org/tools/install>
- Or, update existing install
  - `rustup update`
- Check the environment:
  - `rustup show`



## The Rust Compiler

---

- **rustc** is the main rust compiler
  - Open source, used for multiple targets
  - `rustup target list` to see supported targets
  - `rustup target add <target>` to [add a target support](#)
- Commercial alternatives exist for ISO 26262's ASIL D and IEC 61508's SIL 4 certification levels
  - Hitech (tailored for AURIX™ TC3x and TC4x microcontrollers)
  - AdaCore
  - Ferrocene



## Cross-Compilation with `cargo`

- `Cargo` is both a package manager and build tool
  - Dependencies are pulled from `crates.io`
  - Can be pointed to the `file system` or `github` repo
- The `cargo` command is typically used in favour of `rustc`
  - The `--target` flag is provided to `cargo / rust` for cross compiling

```
cargo build --target thumbv7em-none-eabihf
```

- The default generated output will be using the ELF format
  - Not to be confused with a bin file



## Beyond cargo build: binutils

- Cargo `subcommands` wrapper can use the LLVM tools for working with binaries and object files
  - `nm` lists labels used inside an object
  - `size` provide size information about sections
  - `objdump` disassembler

```
cargo install cargo-binutils
rustup component add llvm-tools
```



## Beyond cargo build: cargo-generate

`cargo-generate` used for creating new cargo projects from [templates](#)

- Augmented version of `cargo new`
- Facilitates the creation of baremetal projects
  - Provides missing linker scripts and projects settings
  - Only need to fill in the missing information (name, target hardware specifics)

```
cargo install cargo-generate
```

```
cargo generate --git https://github.com/rust-embedded/cortex-m-quickstart
cargo generate esp-rs/esp-idf-template cargo
```



## Cargo Runners

Runners are **configuration options** specifying a command to run your binary instead of executing it directly.

- Useful for cross-compilation scenarios (targets, debuggers, emulators..)
- Located inside: `.cargo/config.toml`

Running GDB for a Arm Cortex-M target:

```
[target.'cfg(all(target_arch = "arm", target_os = "none"))']  
runner = "arm-none-eabi-gdb -q -x openocd.gdb"
```



## The Embedded Software Developer's Tool Box

Traditional software development tools can be used with Rust binaries.

- Querying file:
  - `file exec.elf`
- Looking for strings in a file
  - `strings exec.elf`
- Reading ELF file information
  - `readelf -h exec.elf`
  - `readelf --segments exec.elf`
- Hex file dump:
  - `hexdump -C exec.elf`
  - `xxd exec.elf | less`
- Disassembling:
  - `objdump -d exec.elf`



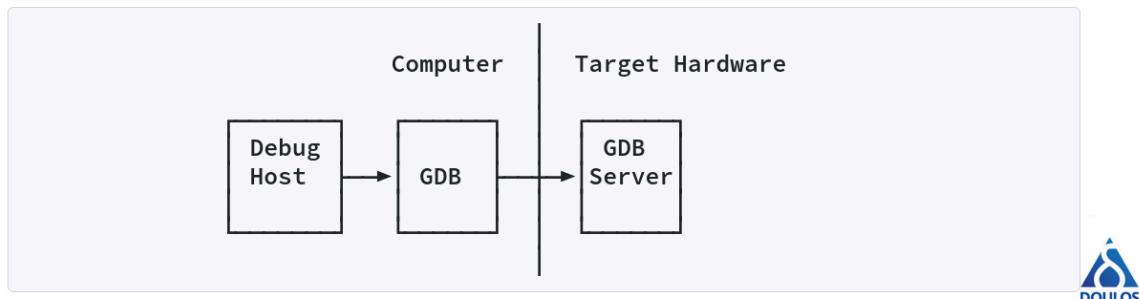
---

# Interactive Debugging



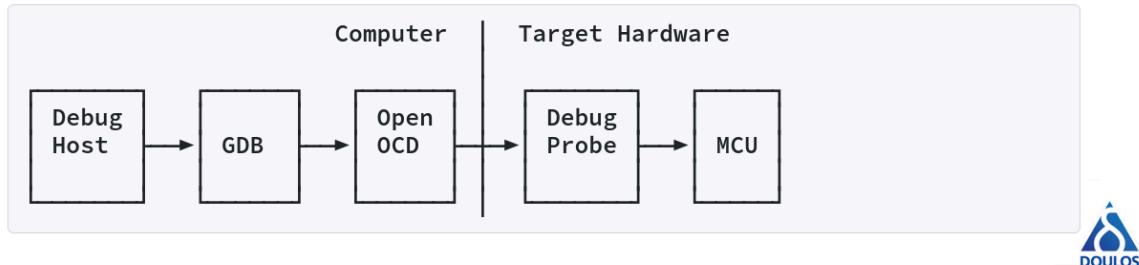
## GDB Overview

- The GNU Debugger is used to debug applications including Rust
- Command line driven (`step`, `continue`, `break ..`)
- Can be used with:
  - `cargo-embed`, through `probe.rs` (simpler)
  - `OpenOCD` more traditional and generic



## OpenOCD Overview

- Open On-Chip Debugger interfaces GDB with a target debug adapter
  - Provide the GDB server required by GDB
  - Supports JTAG and SWD interfaces
- Some hardware solution bypass the need for OpenOCD:
  - Black magic probe (provides a GDB Server)



## Basic Printing Through Semihosting

Semihosting is a mechanism that lets embedded devices do **I/O on the host** and is mainly used to **log messages** to the host console.

- Minimal requirements (a debug session)
  - Lowest common denominator across most MCUs
  - Convenient to use
  - Very slow (hundreds of milliseconds)
- For Arm cortex-m MCUs use: `cortex-m-semihosting` crate
  - You do need to enable semihosting in OpenOCD from GDB first:

```
(gdb) monitor arm semihosting enable
```



## Semihosted Logging Example

- `hprintln!()` uses the same type of formatting as for `println!()` can be used
- The `dbg!()` macro can also be used for quick debugging

```
#![no_main]
#![no_std]

use panic_halt as _;
use cortex_m_rt::entry;
use cortex_m_semihosting::hprintln;

#[entry]
fn main() -> ! {
    hprintln!("Hello, world!").unwrap();
    loop {}
}
```







## Project Configuration and Hardware Platforms

Rust for firmware development

## Aims and Topics

---

In this module we will learn about the overall structure of a embedded-Rust project as well as what constitute a `no_std` rust program.

- Project's structure
- The `no_std` embedded development environment
- Commonly used embedded crates



# Project's Structure



## Project's Files Overview

- `Cargo.toml` project configuration information, dependencies specification
- `config.toml` build information, run commands short cuts (flash/debug..)
- `memory.x` is a linker script short-hand
- Optionally, `build.rs`, `Embed.toml`

```
.  
|-- .cargo  
|   |-- config.toml  
|   |-- Cargo.lock  
|   |-- Cargo.toml  
|   |-- memory.x  
|   |-- build.rs [OPTIONAL]  
|   |-- Embed.toml [OPTIONAL]  
|   |-- src  
|     |-- main.rs
```



## Minimal Cargo.toml File

- `profile.dev`
  - Optimisation level set to 0
  - Debug info on
- `profile.release`
  - Optimisation level 3
  - Debug info off
- Panic strategy has to be defined for an embedded application

```
[package]
name = "crate_name"
version = "0.1.0"
authors = ["Author Name <author@example.com>"]

# the profile used for `cargo build`
[profile.dev]
panic = "abort" # disable stack unwinding on panic

# the profile used for `cargo build --release`
[profile.release]
panic = "abort" # disable stack unwinding on panic
```



## Customising Profiles Inside `Cargo.toml`

Profiles can be customized extensively.

- Optimisation level (`opt-level`)
  - `0..3`
  - Image's size: `s`, `z`
- Debug symbols (`debug`)
  - `true` / `false`
- Link-time optimisation (`lto`)
  - `true` / `false`
- Panic (`panic`)
  - `unwind` / `abort` (better for embedded)

```
[profile.dev]
opt-level = "0"
debug = true
panic = "abort"
```

```
[profile.release]
opt-level = "s"
debug = false
panic = "abort"
lto = true
```



## Adding Project's Dependencies to Cargo.toml

```
[dependencies]
# Regular dependencies for the Cortex-M architecture
cortex-m = "0.7"
cortex-m-rt = { version = "0.7", features = ["device"] }
panic-halt = "0.2"

# A hardware abstraction layer for a specific microcontroller family
stm32fixx-hal = { version = "0.6", features = ["rt", "stm32f103", "medium"] }

# Use no default features to minimize the footprint
embedded-hal = { version = "0.2.4", default-features = false }

# Conditional dependency only for ARM architecture without an OS
[target.'cfg(all(target_arch = "arm", target_os = "none"))'.dependencies]
cortex-m-log = { version = "0.8", features = ["itm"] }

[dev-dependencies]
# Dependencies used only for compiling and running tests, not included in the release builds
proptest = "0.10.1"
```



## Cargo: The `config.toml` File

File used in the context of [embedded software](#) development.

- Defines flags for the [compiler](#)
- [Linker](#) arguments
  - Scripts, optimisations..
  - Additional debug features like defmt..
- [Build target](#) architecture for the Cargo build command
- [Runners](#) specifying how to run your application
  - Flash it
  - Run it on QEMU with GDB...



## config.toml File Example

```
[target.thumbv7m-none-eabi]
runner = "qemu-system-arm -cpu cortex-m3 -machine lm3s6965evel -nographic -semihosting-config enable=on,target=native -kernel"

rustflags = [
    "-C", "linker=flip-link",
    "-C", "link-arg=-Wl,--nmagic",
    "-C", "link-arg=-Wl,-Tlink.x",
    "-C", "link-arg=-Wl,-Tdefmt.x",
    "-C", "inline-threshold=5",
    "-C", "no-vectorize-loops",
]

[build]
# target = "thumbv6m-none-eabi"          # Cortex-M0 and Cortex-M0+
target = "thumbv7m-none-eabi"           # Cortex-M3
# target = "thumbv7em-none-eabihf"        # Cortex-M4F and Cortex-M7F (with FPU)
# target = "thumbv8m.base-none-eabi"       # Cortex-M23
# target = "thumbv8m.main-none-eabi"        # Cortex-M33 (no FPU)
```



## Cargo: The `memory.x` Linker Script

File used to describe the `memory layout` of the target device.

- This includes defining the sizes and boundaries of different memory regions
- Used by crates like `cortex-m-rt`, `esp32c_rt`, etc.
- The `link.x` cross references the `memory.x` definitions

```
MEMORY
{
    FLASH (rx) : ORIGIN = 0x08000000, LENGTH = 256K
    RAM (rwx)  : ORIGIN = 0x20000000, LENGTH = 40K
}
```



---

# The no\_std Embedded Rust Development



## What Does `#![no_std]` Mean?

`#![no_std]` is a **crate level attribute** found inside the `main.rs` or `lib.rs` file.

- Indicates that the crate will link to the `core` crate instead of the `std` crate
- The `core` crate is a **subset** of the `std` crate that makes zero assumptions about the system the program will run on
  - Lacks the APIs for anything that involves heap memory allocations, networking and I/O
  - Must initialize its own runtime



## How Are Attributes Used?

Rust has a rich set of attributes that can be used to **control** various aspects of code **compilation** and **behaviour**.

- Crate-Level: `#![no_std]`, `#![feature]`, `#![no_main]`
- Conditional Compilation: `#[cfg]`, `#[cfg_attr]`
- Lint: `#[allow]`, `#[warn]`, `#[deny]`
- Derive: `#[derive]`
- Function and Method: `#[inline]`, `#[test]`
- Item: `#[repr]`
- Many more...



## Module and Item Attributes Example

Attributes that apply to modules, structs, enums, and other items.

`#[repr(...)]` : Controls the memory layout of structs and enums.

```
#[repr(C)]
struct MyStruct {
    field1: i32,
    field2: u8,
}
```

`#[repr(align(N))]` : Used to control the memory layout of structs and enums

```
#[repr(align(16))]
struct AlignedStruct {
    data: u32,
    more_data: u64,
}
```



## Minimal Freestanding Rust Binary

```
#![no_std] // don't link the Rust standard library
#![no_main] // disable all Rust-level entry points

use core::panic::PanicInfo;

#[no_mangle] // don't mangle the name of this function
pub extern "C" fn _start() -> ! {
    // this function is the entry point, since the linker looks for a function
    // named `_start` by default
    loop {}
}

/// This function is called on panic.
#[panic_handler]
fn panic(_info: &PanicInfo) -> ! {
    loop {}
}
```





## Processors' Support

---

A **wealth** of processor architectures are **supported** by the Rust ecosystem either at the PAC level or PAC/HAL level.

- Arm based MCU/CPU
  - Infineon, NXP, STMicroelectronics, Nordic, Atmel, Renesas
- RISC-V based processors
- Infineon TriCore
- Expressif ESP32 processors
- Microchip AVR
- Texas Instruments MSP
- MIPS



## Arm Cortex-\* Crates

- Cortex-A/R/M are at the core of many SoCs
- Cortex-M is the **dominant** generic MCU architecture
- Some relevant Cortex-M crates:
  - `cortex-m` : PAC for the core peripherals (timer, debugger, etc..)
  - `cortex-m-rt` : minimal runtime (entry point, interrupts, exception, ...)
  - `panic-itm` : Panic handler that sends messages over the ITM/SWO output
  - `panic-semihosting` : Panic handler that sends messages over semihosting



## The Run-time cortex-m-rt Crate

- Specifies how does our code gets **located** inside our MCU **memory map**
- Linker script (generic for Cortex-m)
- Locates the vector table
- Initialise static variables
- Enable FPU (if it exists)
- Provides function attributes
  - `#[entry]` to declare the **entry point** of the program
  - `#[exception]` to override an exception handler
    - All exceptions default to an infinite loop
  - `#[pre_init]` to run code before static variables are initialised



## cortex-m-rt Functions' Attributes Example

```
// IMPORTANT the standard `main` interface is not used because it requires nightly
#![no_main]
#![no_std]

// Some panic handler needs to be included. This one halts the processor on panic.
use panic_halt as _;

use cortex_m_rt::entry;

// use `main` as the entry point of this application
// `main` is not allowed to return
#[entry]
fn main() -> ! {
    // initialization
    loop {
        // application logic
    }
}
```



## Cortex-M Micro-Architectural Crate Example

The `SysTick` peripheral that's common to all Cortex-M based MCUs.

```
#![no_std]
#![no_main]

use cortex_m::peripheral::{syst, Peripherals}; // ←💡 Cortex-M micro-architectural crate
use cortex_m_rt::entry;
use panic_halt as _;

#[entry]
fn main() -> ! {
    let peripherals = Peripherals::take().unwrap();
    let mut systick = peripherals.SYST; // ←💡 Timer common to all Cortex-M
    systick.set_clock_source(syst::SystClkSource::Core);
    systick.set_reload(1_000);
    systick.clear_current();
    systick.enable_counter();
    while !systick.has_wrapped() {}
    loop {}
}
```







## Embedded Software Development

Rust for firmware development

## Aims and Topics

---

In this chapter we will learn about the specific aspects related to software development for an embedded application.

- Embedded Rust ecosystem
- Mixing Rust with assembly

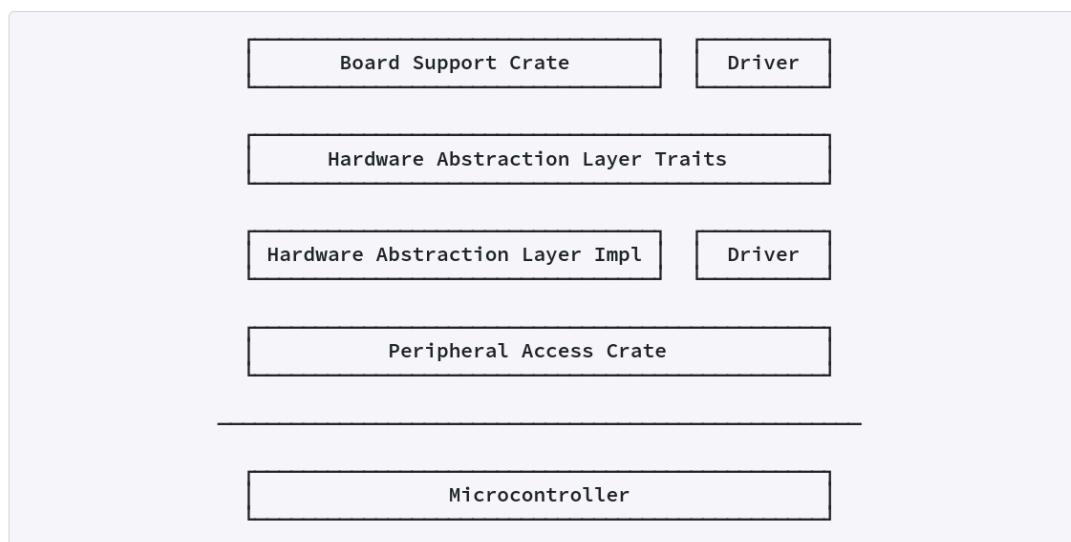


---

# Embedded Rust Ecosystem



## Abstraction Layers



## Unsafe Rust

---

**Memory** locations of **mapped** hardware can be accessed just like in embedded C but it is **unsafe**

- Large amount of technical knowledge needed when controlling hardware
  - Read the (very large) manual!
- Very error prone
  - No safety net, no validity or dependencies checks



## Unsafe Code Example

```
static HELLO: &[u8] = b"Hello World!";

#[entry]
fn main() -> ! {
    let vga_buffer = 0xb8000 as *mut u8; // raw pointer to memory

    for (i, &byte) in HELLO.iter().enumerate() {
        unsafe { // unsafe block for dereferencing the raw pointers
            *vga_buffer.offset(i as isize * 2) = byte;
            *vga_buffer.offset(i as isize * 2 + 1) = 0xb;
        }
    }

    loop {}
}
```



## Peripheral Access Crates

PACs provide low **level access** to memory mapped peripherals inside a given MCU

- Adds **zero cost abstraction**
- Usually auto-generated by the **SVD2Rust** program
  - Using the CMIS-SVD format
- Limited usefulness due to MCU complexity
  - Maximum control of the hardware
  - Stepping stone to the creation of an HAL



## The PAC Application Programming Interface

A **singleton** design pattern is used for peripherals

- Ensures only one instance of device peripherals can be used

```
let core_peripherals = stm32f401_pac::Peripherals::take().unwrap();
```

- `take()` method returns an `Option<T>`.
  - First call returns an instance; Subsequent calls return `None`.
- The `core_peripheral` handle is used to create instances for specific peripherals.



## Peripheral Access Manipulation

- Access format
  - **read** : Read values from device registers.
  - **modify** : Modify specific bits in device registers.
  - **write** : Write new values to device registers.
- Usage Format:

[Peripheral Handle].[Peripheral Register Name].[Operation]

```
let dp = stm32f401_pac::Peripherals::take().unwrap();
// Enabling the clocks for GPIOA and GPIOC
dp.RCC.write(|w| w.gpioaen().set_bit().gpiocen().set_bit());
// Toggling PA6 Output Using the modify Method
dp.GPIOA.odr.modify(|r, w| w.odr6().bit(!r.odr6().bit()));
```



## PAC Use Example

```
#![no_main]
#![no_std]

use panic_halt as _;

use cortex_m_rt::entry;
use nrf52833_pac::Peripherals;

#[entry]
fn main() -> ! {
    if let Some(p) = Peripherals::take() {
        p.P0.pin_cnf[28].write(|w| w.dir().output()); // ← This configures pins 28 and 21 of port P0 as output pins
        p.P0.pin_cnf[21].write(|w| w.dir().output()); // ←

        p.P0.out.write(|w| unsafe { w.bits(1 << 21) }); // ← Setting the initial state of pin 21 to high

        let mut count: u8 = 0;
        loop {
            count += 1;

            if count & 1 == 1 {
                p.P0.out.write(|w| unsafe { w.bits(1 << 21) }); // ← Setting the pin 21 to high
            } else {
                p.P0.out.write(|w| unsafe { w.bits(0) }); // ← Setting the pin 21 to low
            }

            for _ in 0..50_000 {
                cortex_m::asm::nop();
            }
        };
    }

    loop {
        continue;
    }
}
```



## What is an HAL?

---

An Hardware Abstraction Layer provides **portability** for embedded systems with varying peripherals capabilities across different vendors and families

- Uniform interface for different hardware platforms through standardisation
- Developers can write device independent applications
- Abstraction of platform-specific details
- Active community and ecosystem support
  - [crates.io](https://crates.io)
  - HAL implementation crates on <https://github.com/rust-embedded/awesome-embedded-rust>



## HAL Example

```

#[no_main]
#[no_std]

use panic_halt as _;

use embedded_hal::digital::OutputPin; // Provides required traits
use cortex_m_rt::entry;
use nrf52833_hal as hal;
use nrf52833_hal::{gpio::Level, pac::Peripherals};

#[entry]
fn main() -> ! {
    if let Some(p) = Peripherals::take() {
        let port0 = hal::gpio::p0::Parts::new(p.P0); // ← Split the GPIO port P0 into individual pins for configuration

        // Configure P0.28 and P0.21 as output pins
        let mut col1 = port0.p0_28.into_push_pull_output(Level::Low); // ← Configure P0.28 as a push-pull output, initially set to Low
        let mut row1 = port0.p0_21.into_push_pull_output(Level::High); // ← Configure P0.21 as a push-pull output, initially set to High
        let mut count: u8 = 0;
        loop {
            count += 1;

            if count & 1 == 1 {
                row1.set_high().unwrap(); // ← Set P0.21 (row1) to High (turn on the corresponding output)
            } else {
                row1.set_low().unwrap(); // ← Set P0.21 (row1) to Low (turn off the corresponding output)
            }

            for _ in 0..50_000 {
                cortex_m::asm::nop();
            }
        }
    }
}

```



## Embedded Rust Board Support Package

The BSPs are an addition to HALs as they abstract board specific feature for given boards

- Defines specific pins, MCU configurations, hardware for a board

```
#![no_main]
#![no_std]

use panic_halt as _;
use cortex_m_rt::entry;
use embedded_hal::{delay::DelayNs,
                   digital::OutputPin};
use microbit::{board::Board,
               hal::timer::Timer};

#[entry]
fn main() -> ! {
    let mut board = Board::take().unwrap();
    let mut timer = Timer::new(board.TIMER0);
    let _ = board.display_pins.col1.set_low();
    let mut row1 = board.display_pins.row1;
```

```
// continuing...
loop {
    let _ = row1.set_low();
    timer.delay_ms(1_000u32);
    let _ = row1.set_high();
    timer.delay_ms(1_000u32);
}
```



---

# Mixing Rust with Assembly



## Why Use Intrinsics?

- Intrinsics provide access to specific CPU **instructions** that are **not** typically **exposed** by higher-level languages
- Useful for **performance optimisation** and accessing specialised hardware features
  - **Limited** scope: `wfe`, `wfi`, `dsb`, `isb`, `dmb`, `bkpt` ...

```
use cortex_m::asm::isb;
fn main() {
    // Use the Instruction Synchronization Barrier
    unsafe {
        isb();
    }
}
```



## Inline Assembly

Inline assembly is provided by the macro `asm!()`

- Can be used to embed handwritten assembly in the output generated by `rustc`
  - x86, x86-64, AArch64, AArch32, RISC-V, LoongArch, TriCore...
- Inference example:

```
fn reverse_bytes_manual(x: u32) -> u32 {  
    ((x & 0x000000FF) << 24) |  
    (((x & 0x0000FF00) << 8) |  
     ((x & 0x00FF0000) >> 8) |  
     ((x & 0xFF000000) >> 24)  
}
```

- Inline assembly example:

```
fn reverse_bytes(x: u32) -> u32 {  
    let result: u32;  
    unsafe {  
        asm!(  
            "rev {0}, {1}",  
            out(reg) result, // Output operand  
            in(reg) x // Input operand  
        );  
        result  
    }  
}
```



## Module Level Inline Assembly

The `global_asm!()` macro is used for defining **startup code**, interrupt **vector tables**, or other essential assembly routines that are too **large** or complex for `asm!()`.

- Functions can be written completely in assembly
  - Used for startup routines that initialise hardware and memory before Rust's main function starts
  - Used for defining an interrupt vector table

```
global_asm!(r#"
.globl my_assembly_func

my_assembly_func:
    bx lr
"#);

extern {
    fn my_assembly_func();
}
```







## Mixing Rust with C

Rust for firmware development

## Aims and Topics

---

In this module we will look at the foreign functions' interface and how to use it to mix both C and Rust code.

- FFI and ABI in Rust
- Calling C from Rust
- Calling Rust from C





## What is a Foreign Function Interface?

---

FFI is a system-level mechanism for **interacting** with different programming **languages**.

- In the Rust context we want to:
  - Facilitate the **reuse** of existing **libraries** written in languages like C within Rust code
  - **Expose Rust** functions for use in C or C++ projects, enhancing safety through Rust's guarantees



## Procedure Call Standards in Rust, C, and Assembly

- The PCS is a component of the ABI that focuses on [function calls](#)
  - Rust relies on the ABI defined by the target platform (AArch64, AArch32, etc.)
  - Pragmatically, Rust commonly uses the C calling conventions
- [AArch32 PCS Example:](#)
  - First 4 function arguments map to R0-R3
  - Additional arguments are stacked
  - 64-bit types map to consecutive registers (even index first)
  - The stack pointer is 8-byte aligned at the object's boundaries



## C ABI: Register Usage Example in Calls

---



## C ABI: Data Alignments of Arrays

All values have an **alignment** and **size**.

- The alignment of a value specifies which **addresses** are **valid** for storing the value.
  - For example, a `u32` has a size of 4 bytes and an alignment of 4 (`size_of::<u32>()`).
- Proper alignment ensures **efficient** data access and can **prevent hardware exceptions** on some architectures.
- An array of `[T; N]` (a fixed-size array in Rust) has a size of `size_of::<T>() * N` and the same alignment as `T`



## C ABI: Data Alignments of Structures

The alignment of a `struct` with a C representation is the alignment of the most-aligned field in it

```
use std::mem::{size_of, align_of};

#[repr(C)]
struct ExampleStruct {
    field1: u8,           // 1-byte alignment
    field2: u32,          // 4-byte alignment
    field3: i16,          // 2-byte alignment
}

fn main() {
    // Assertions for size and alignment checks
    assert_eq!(size_of::<ExampleStruct>(), 12);
    assert_eq!(align_of::<ExampleStruct>(), 4);
}
```

```
use std::mem::{size_of, align_of};

#[repr(C)]
struct ExampleStruct {
    field_1: u8,           // 1-byte alignment
    padding_1: [u8; 3],    // Padding for alignment
    field_2: u32,          // 4-byte alignment
    field_3: i16,          // 2-byte alignment
    padding_2: [u8; 2],    // Padding for alignment
}

fn main() {
    // Assertions for size and alignment checks
    assert_eq!(size_of::<ExampleStruct>(), 12);
    assert_eq!(align_of::<ExampleStruct>(), 4);
}
```



## C ABI: Data Alignments of Enums

The use of `#[repr(C)]` to ensure that enums have a predictable memory layout compatible with C.

- The Enum size determination is based on the largest variant and the discriminant

```
use std::mem::{size_of, align_of};
#[repr(C)]
enum Status {
    Idle, // No associated data
    Active(u32), // Single u32 data
    Error { code: u16, desc: u16 }, // Struct-like variant with two u16 fields
}

fn main() {
    println!("Size of Status enum: {}", size_of::<Status>());
    // Expected: 8 bytes (4 bytes discriminant + 4 bytes largest variant)

    println!("Alignment of Status enum: {}", align_of::<Status>());
    // Expected: 4 bytes, matching the strictest field (u32)
}
```



---

# Calling C from Rust



## Naming Convention Compatibility

Rust has well-defined [naming conventions](#) to ensure code consistency and safety.

- Compiler generates [warnings](#) if Rust naming conventions are violated
- C/C++ do not enforce strict naming conventions
- Crate attributes are required to [suppress warnings](#) when interfacing with C/C++ code

```
#![allow(non_upper_case_globals)]
#![allow(non_camel_case_types)]
#![allow(non_snake_case)]
```



## Declaring External Functions in Rust

The `extern` keyword is used to declare functions that are implemented in other languages.

- Alternatively, use `extern "aapcs"` for ARM-specific calling conventions

```
extern "C" {
    fn SysTick_Handler();
}
```

- This assumes the existence of the C function with the following definition:

```
void SysTick_Handler(void) /* some code */
```



## Handling Scalar Data Types

The `core::ffi` module in Rust, provides various type definitions for C types.

C type	core::ffi equivalent	C type	core::ffi equivalent
char	c_char	long	c_long
signed char	c_schar	unsigned long	c_ulong
unsigned char	c_uchar	long long	c_longlong
int	c_int	unsigned long long	c_ulonglong
unsigned int	c_uint	double	c_double
short	c_short	float	c_float
unsigned short	c_ushort	void pointer	c_void



## Example of Scalar Types Usage

```
#![allow(non_upper_case_globals)]
#![allow(non_camel_case_types)]
#![allow(non_snake_case)]

use core::ffi::{c_char, c_int};

extern "C" {
    fn microbit_display_scroll(s: *const c_char);
    fn wait_ms(s: c_int);
}

fn main() {
    unsafe {
        // Calling the external C function within an unsafe block
        microbit_display_scroll("Hello Rust!\0".as_ptr());
        wait_ms(2000);
    }
}
```



## Direct and Indirect Linking

- **Direct Linking:** Ensures that your Rust application links directly to a specific C library using the `#[link]` attribute.
- **Indirect Linking:** Managed by build environments like Cargo, where linking is specified in a `build.rs` file or through Cargo configurations.

```
// Direct linking
#[link(name = "microbit_Screen")]

extern "C" {
    fn microbit_display_scroll(s: *const c_char);
    fn wait_ms(s: c_int);
}
```



## Working with C enum and struct types

1. Define the enum or stuct in C
2. Define the matching version in Rust with the `#[repr(C)]` attribute

```
enum IconNames {
    Heart = 0,
    SmallHeart = 1,
    Yes
};
```

```
#[repr(C)]
pub enum IconNames {
    Heart = 0,
    SmallHeart = 1,
    Yes
}

extern "C" {
    fn showIcon(icon: IconNames, interval: c_uint);
}
```

```
typedef struct TPoint { uint8_t x; uint8_t y; };
```

```
#[repr(C)]
pub struct TPoint {
    pub x: u8,
    pub y: u8,
}

extern "C" {
    // External C function that accepts a TPoint
    fn update_led_position(point: TPoint);
}

fn main() {
    let mut point = TPoint { x: 10, y: 20 };

    // C function that updates the LED's position
    unsafe {
        update_led_position(point);
    }
}
```



---

# Calling Rust from C



## Rust's Library Specification...

Rust functions are placed in the `lib.rs` file

- `#[no_mangle]` : **Avoids** the Rust compiler symbols' **mangling**
- `extern "C"` : Makes functions adhere to the C calling **convention**

```
#![no_std]

use core::panic::PanicInfo;

#[panic_handler]
fn handle_panic(_info: &PanicInfo) -> ! { loop {} }

#[no_mangle]
pub extern "C" fn rust_callee(a: i32, b: i32, c: i32, d: i32, e: i32, f: i32) -> i32 {
    a + b + c + d + e + f
}
```



## Rust's Library Specification

The compiler needs additional library definition directives inside the `Cargo.toml` file

- `crate-type` : Determines the type of library to generate, such as static or dynamic libraries, to match the linking requirements
  - Add to the Cargo.toml file the following directive:

```
[lib]
name = "rust_lib"
crate-type = ["staticlib"] # static library (.a)
#           ["cdylib"]      dynamic library (.so)
```



## On the C Side

- Declare the Rust function in a C-compatible header file for proper linkage
- Use the `extern` keyword to reference Rust functions within C source files

```
// rust_lib.h
int rust_callee(int a, int b, int c, int d, int e, int f);
```

```
// main.c
extern int rust_callee(int a, int b, int c, int d, int e, int f);

int main(void) {
    rust_callee(1, 2, 3, 4, 5, 6);
    return 0;
}
```



## Compiling and Linking

Assuming the C file is place at the crate's root. We can compile is as shown:

- **Linker Arguments `-l` and `-L`**: Indicate the library name and the directory containing the Rust object files
- **Setting `LD_LIBRARY_PATH`**: Specifies the runtime library search path for dynamic libraries

```
gcc call_rust.c -o call_rust -lrust_lib -L./target/debug  
LD_LIBRARY_PATH=./target/debug ./call_rust
```







## Memory Management in no\_std

Rust for firmware development

## Aims and Topics

---

In this module we will look at the various kinds of memory used by a embedded-Rust program and look at pragmatic options for using complex types.

- Stack and heap memory allocation
- Custom allocators for heap usage
- Alternative to using an allocator



---

# Stack/Heap Memory Allocations



## Why Should I Care About It?

---

- Efficient Resource Utilisation in Embedded Rust:
  - Limited memory in microcontrollers (MCUs)
  - Proper management prevents memory exhaustion
- Real-time Constraints:
  - Memory management consumes CPU time
- Safety and Security:
  - Prevents overflows and ensures proper memory access



## Memory in an Embedded Program

	Content	Size	Lifetime	Cleanup
Stack ↓	+Function arguments +Local variables +Known size at compile time	Dynamic / fixed upper limit	Lifetime of function	Automatic. When function returns
↑ Heap	+Values that live beyond a function's lifetime +Values access by multiple threads +Large values +Unknown size at compile time	Dynamic	Determined by programmer	Manual Malloc/ free MMU/MPU Operations
Static	+Program' binary +Static variables +String literals	Fixed	Lifetime of program	Automatic. When program terminates



## Rust Use of Stack Memory Resources

---

- Efficient and quick memory for local variables:
  - Often cached in embedded systems
- Local variables in functions:
  - Local variables declared in a function are added to that function's stack frame
- Type restrictions:
  - Only values with a size known at compile time (e.g., `i8`, `u32`) can exist on the stack



## Rust Use of Heap Memory Resources

- Heap Resources:
  - Used for values with variable sizes at **runtime** (e.g., a vector of characters)
  - **Elements** such as characters can be **added** or **removed** during runtime

```
fn main() { //                                └─ (HEAP)
    //                                     ↓
    let a = String::from("Ferris");
} //                                     ↑
//                                     └─ Fat pointer [ptr, len, capacity]
//                                         (STACK)
```



---

## Custom Allocators for Heap Usage



## Rationale for Custom Allocators

---

- Custom allocators can provide essential features in embedded systems:
  - Control over **memory usage** and allocation patterns
  - Real-time **predictability**
  - Eliminates the need for an operating system
  - Targets specialised memory types
- Examples: `embedded_alloc`, `buddy_system_allocator`, `talc`,  
`alloc_no_std_lib`



## The `embedded-alloc` custom allocator

- MCU-agnostic heap allocator for embedded systems
  - Requires an implementation of the `critical-section` feature
- Linked-List First Fit (LLFF) heap allocator implementation:
  - Handles `varying` block `sizes` effectively
  - Performance may degrade with heavy heap usage



## Critical Section Implementations

- A critical section is a fundamental primitive for controlling concurrency
- Ensures mutually exclusive access to a resource
  - Can be acquired by only one thread at a time
- Applicable to various implementations (unicore, multicore, bare metal, RTOS...)

```
[dependencies]
cortex-m = { version = "0.7.7", features = ["critical-section-single-core"] }
embedded-alloc = "0.5.1"
```



## embedded-alloc Example

```
#![no_main]
#![no_std]

extern crate alloc;
use alloc::boxed::Box, vec::Vec;
use cortex_m_rt::entry;
use stm32f1xx_hal;
use embedded_alloc::Heap;
use panic_halt as _;

#[global_allocator]
static HEAP: Heap = Heap::empty();

#[entry]
fn main() -> ! {
    {
        use core::mem::MaybeUninit;
        const HEAP_SIZE: usize = 1024; // 1KiB
        static mut HEAP_MEM: [MaybeUninit<u8>; HEAP_SIZE] = [MaybeUninit::uninit(); HEAP_SIZE];
        unsafe { HEAP.init(HEAP_MEM.as_ptr() as usize, HEAP_SIZE) }
    }

    let mut vec = Vec::new();
    vec.push(true);
    vec.push(false);

    let my_box = Box::new(42);

    #[allow(clippy::empty_loop)]
    loop {}
}
```



---

## Alternatives to Using an Allocator



## The `heapless` Crate

- Provides `fixed-capacity` vectors and strings
- More versatile than `arrayVec` :
  - Implements a map and set based on a hash table
  - Includes more complex structures such as:
    - `Box` , `Arc` , `BinaryHeap` , queues, deques, maps, and more



## heapless::Vec Example

```
use heapless::Vec;

let mut my_vec: Vec<u32, 16> = Vec::new();
let _ = my_vec.push(1963);
let _ = my_vec.push(9);

// Try to add elements to the my_vec
if my_vec.push(1986).is_err() {
    writeln!(serial, "Failed to add element to my_vec").unwrap();
}

// Access elements by index
if let Some(value) = my_vec.get(1) {
    writeln!(serial, "Value at index 1: {}", value).unwrap();
} else {
    writeln!(serial, "No value at index 1").unwrap();
}

// Iterate over elements
for value in &my_vec {
    writeln!(serial, "Iterating value: {}", value).unwrap();
}

// Pop an element
if let Some(value) = my_vec.pop() {
    writeln!(serial, "Popped value: {}", value).unwrap();
} else {
    writeln!(serial, "No value to pop").unwrap();
}
```



## heapless::String Example

```
use heapless::String;

let mut my_string: String<6> = String::new();
let _ = my_string.push_str("Ferris");
writeln!(serial, "{} quotes: i don't believe in isims.\r", my_string).unwrap();

let my_string: String<9> = String::from_str("Bueller's").unwrap();
writeln!(serial, "{} quotes: i don't believe in isims.\r", my_string).unwrap();

let mut subject = String::new(); // Create a new String with capacity 9
let _ = subject.push_str("Ferris"); // Push "Ferris" into the String

// Try to append " Bueller's" to the String
if subject.push_str(" Bueller's").is_ok() {
    writeln!(serial, "{} quotes: I don't believe in isms.\r", subject).unwrap();
} else {
    writeln!(serial, "String error").unwrap();
}
```



## Compromise Container Alternatives

- Store a certain number of elements inline
  - Fall back to the heap (entirely) for larger allocations
- Improves cache locality
- Reduces allocator traffic for small workloads
- `Smallvec` or `Tinyvec` crates
  - Requires an `custom allocator` to use the heap in `no_std`

```
use smallvec::SmallVec;  
  
// spills into the heap over 4 elements  
let mut v = SmallVec::<[u8; 4]>::new();
```







## Bringing a Baremetal System to Life

Rust for firmware development

## Aims and Topics

---

In this chapter we will look at the steps involved in starting a main function onto a deeply embedded device.

- Tailoring the image memory map to your target
- Rust library initialisation
- Streamlining bare-metal development
- Exception handlers

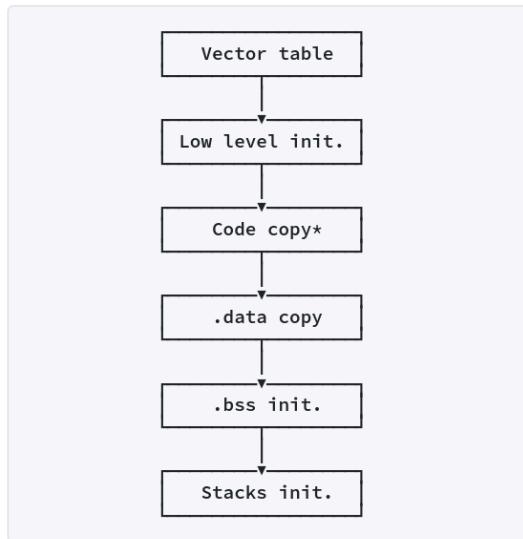


---

# Processor Initialisation



## General boot sequence overview



1. The **vector** table provide the reset handler's address
2. System devices' **initialisation** may be required
3. The code can be **optionally** moved to the RAM
4. **Static data** requires initialisation
5. Optional additional **stacks** can be configured



## The Vector Table: Cortex-M Example

Vector tables will **differ** from one processor's family to another.

- They contain either **jump** instructions or a **destination** address
  - Both internal and external exceptions are listed
- With Arm Cortex-M processors, it also contains the MSP address

0x40 + 4*N	External N
...	...
0x40	External 0
0x3C	SysTick*
0x38	PendSV
0x30 to 0x34	Reserved (x2)
0x2C	SVCall
0x10 to 0x28	Reserved (x7)
0x0C	Hard Fault
0x08	NMI
0x04	Reset
0x00	main stack ptr



## Writing an Exceptions' Table in Rust

The vector table entries can either be a [pointer](#) to an exception handler [or](#) a reserved [placeholder](#).

```
pub union Vector {
    reserved: u32,
    handler: unsafe extern "C" fn(),
}

extern "C" {
    fn NMI();
    fn HardFault();
    fn SVC();
    fn PendSV();
    fn SysTick();
}
```

```
#[link_section = ".vector_table.exceptions"]
#[no_mangle]
pub static EXCEPTIONS: [Vector; 14] = [
    Vector { handler: NMI },
    Vector { handler: HardFault },
    Vector { reserved: 0 },
    Vector { handler: SVC },
    Vector { reserved: 0 },
    Vector { reserved: 0 },
    Vector { handler: PendSV },
    Vector { handler: SysTick },
];
```



---

# Rust Library Initialisation



## General Form of a Reset Handler Function

- The reset handler has some specific memory [placement requirements](#)
- We need to provide [section attributes](#) for the benefit of the linker

```
pub fn reset_handler() -> ! { // <-- returns the never type
    // handler's code here....
    loop {}
}

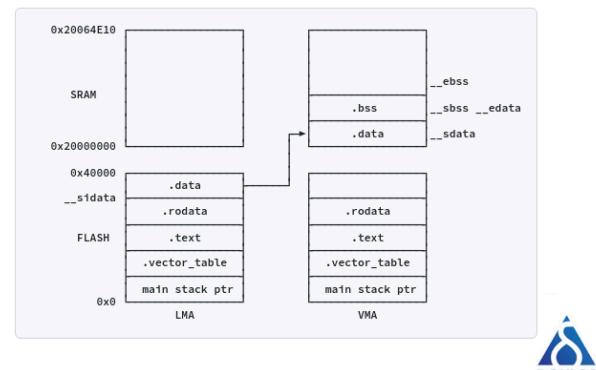
// The reset vector, a pointer into the reset handler
#[link_section = ".vector_table.reset_vector"]
#[no_mangle]
pub static RESET_VECTOR: unsafe extern "C" fn() -> ! = reset_handler;
```



## Providing Symbols for the Reset Handler

- The **initialisation** algorithm relies on the existence of **address symbols** provided by the linker script
  - Those symbols will be used to perform a move operation

```
static mut __sbss: u32; // Start of .bss section
static mut __ebss: u32; // End of .bss section
static mut __sdata: u32; // Start of .data section
static mut __edata: u32; // End of .data section
static __sidata: u32; // LMA of .data section
```



## Basic Reset Handler: Symbols Definition

- The [placement symbols](#) are defined by the [linker](#) script and defined as `extern` values

```
pub fn reset_handler() -> ! { // <- returns the never type
    extern "C" {
        // These symbols come from `linker.ld`
        static mut __sbss: u32; // Start of .bss section
        static mut __ebss: u32; // End of .bss section
        static mut __sdata: u32; // Start of .data section
        static mut __edata: u32; // End of .data section
        static __sidata: u32; // LMA of .data section
    }

    // More code...
    loop {}
}
```



## Basic Reset Handler: R/W Data Section Initialisation

- Data section containing global initialised values

```
pub fn reset_handler() -> ! {
    // linker placement symbols

    // Initialise Data
    unsafe {
        let mut sdata: *mut u32 = &mut __sdata;
        let edata: *mut u32 = &mut __edata;
        let mut sidata: *const u32 = &__sidata;

        while sdata < edata {
            write_volatile(sdata, read(sidata));
            sdata = sdata.offset(1);
            sidata = sidata.offset(1);
        }
    }
    // More code...
}
```



## Basic Reset Handler: Zero Initialised Data Section

- BSS sections only require a zeroing out

```
pub fn reset_handler() -> ! {
    // linker placement symbols

    // Initialise Data

    // Initialise (Zero) BSS
    unsafe {
        let mut sbss: *mut u32 = &mut __sbss;
        let ebss: *mut u32 = &mut __ebss;

        while sbss < ebss {
            write_volatile(sbss, zeroed());
            sbss = sbss.offset(1);
        }
    }

    // Call user's main function
    main()
}
```



## Linker Script with Explicit Address Markers

```
SECTIONS
{
    /* MORE CODE */

    .text :
    {
        __stext = .;
        *(.text .text.*);
        . = ALIGN(4);
        __etext = .;
    } > FLASH

    .rodata :
    {
        __srodata = .;
        *(.rodata .rodata.*);
        . = ALIGN(4);
        __erodata = .;
    } > FLASH
}
```

```
.data :
{
    __sdata = .;
    *(.data .data.*);
    . = ALIGN(4);
    __edata = .;
} RAM AT > FLASH

/* LMA of .data */
__sidata = LOADADDR(.data);

/* .bss */
.bss :
{
    __sbss = .;
    *(.bss .bss.*);
    . = ALIGN(4);
    __ebss = .;
} > RAM
}
```



---

## Steamlining Bare-Metal Development



## Using `#![no_std]` Crate Wide Attribute

- A crate wide attribute is used to `disable` the standard library
- Ensures compatibility with environments lacking an OS

```
#![no_std]

fn main() {
    // Bare metal code here
}
```



## The Smallest `#![no_std]` Program

- The `#![no_main]` attribute which means that the program **won't use** the standard `main` function as its **entry point**
  - This code produces an empty binary!

```
#![no_main]
#![no_std]

use core::panic::PanicInfo;

#[panic_handler]
fn panic(_panic: &PanicInfo<'_>) -> ! {
    loop {}
}
```



## The Cortex-M Runtime (rt) Crate

- Contains all the required parts to build a `no_std` application
  - Vector table configuration
  - Static variables' initialisation
  - Enabling of the FPU
- Specialised attribute definition
  - `#[entry]` program's entry point definition
  - `#[exception]` overriding of existing exception handlers
  - `#[pre-init]` custom code to execute before variable initialisation



## Pragmatic, Minimal Bare-Metal Example

Pragmatically, we do not want to have to re-invent the wheel.

- "rt" core crates provide **attributes** and boiler-plate **initialisation code** to simplify the creation of an application's run-time

```
#![no_std]
#![no_main]

use panic_halt as _;
use cortex_m_rt::entry;

#[entry]
fn main() -> ! {
    loop {}
}
```



## Project's Cargo File Configuration

- The `Cargo.toml` file will lists your project's dependencies
- Use [crates.io](#) to find a suitable crate

```
cortex-m = "0.6.0"
cortex-m-rt = "0.6.10"
stm32l4xx-hal = {version="*", features=["stm32l4x2", "rt"] }
```



---

## Exception Handlers



## Interrupts on MCUs

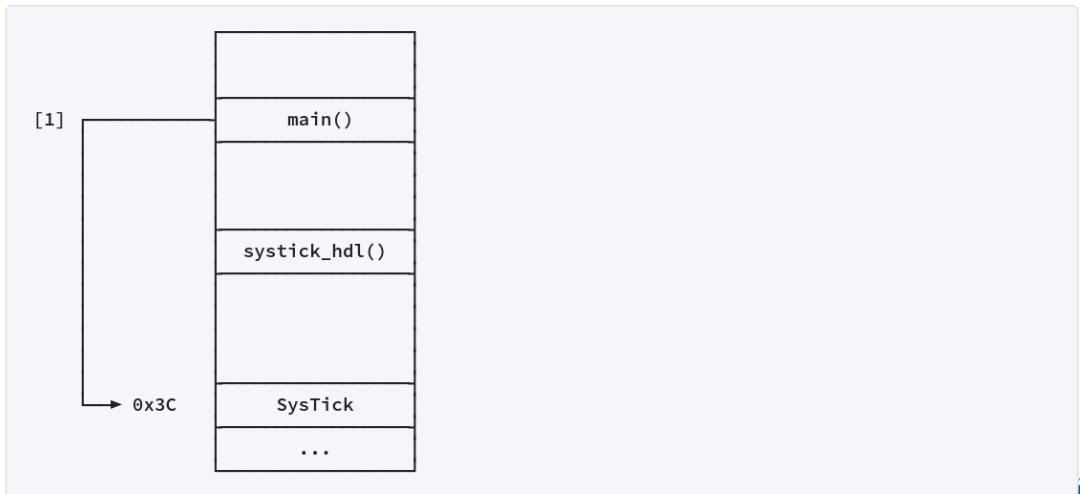
---

- Interrupts are **device specific**
  - Number
  - Source
  - Priority
- Generic requirements:
  - Setup the peripheral for generating interrupts
  - Setup the **priority level** to enable pre-emption
  - **Enable** the source and the **interrupt** controller



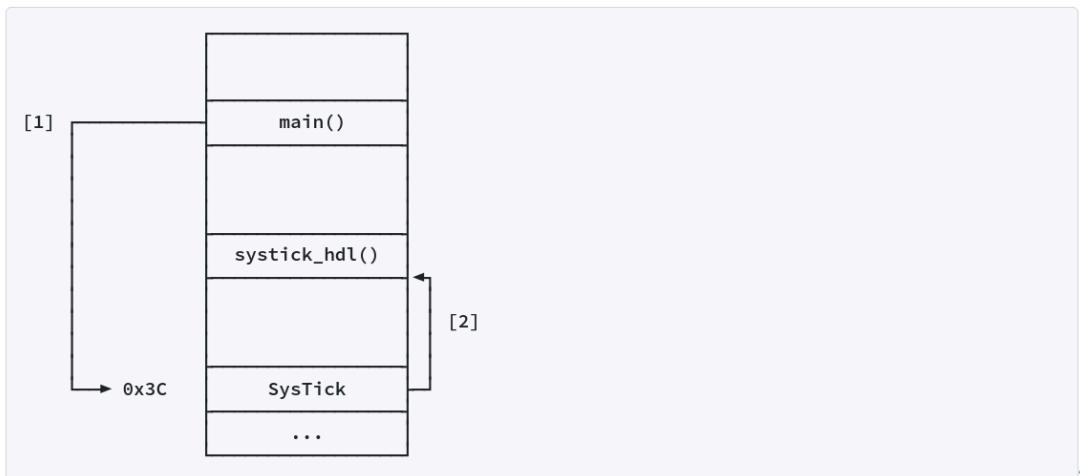
## SysTick Exception Example

- On an exception the core vectors into the relevant entry inside the vector table



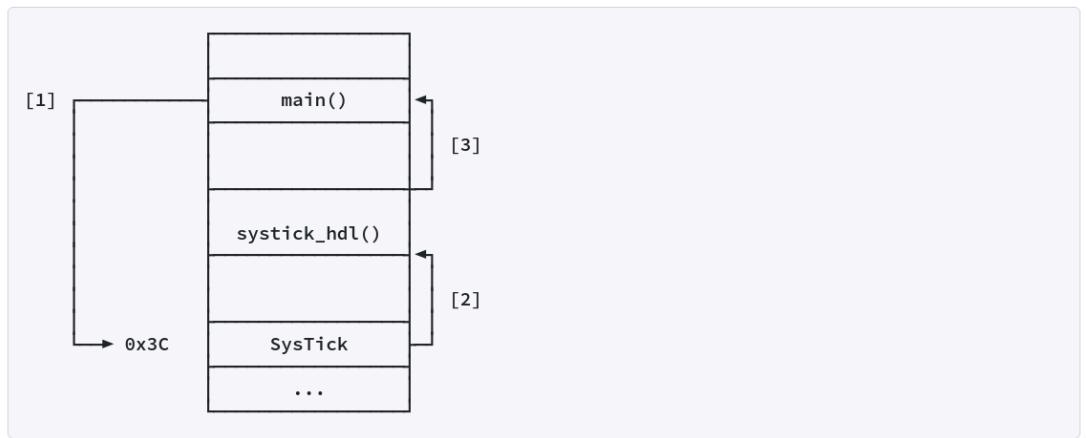
## Jumping to the Exception Handler

- The address held inside the vector table is loaded into the Program Counter



## Execution and Return

- The specific handler is executed and the processor returns the main application



## Handlers with the `cortex_m_rt` Crate

- Handlers are defined by overriding default implementations
  - Use the semantics: `#[exception] fn Name(..)`
- Supported handlers' name:

```
DefaultHandler
NonMaskableInt
HardFault
MemoryManagement(*)
BusFault(*)
UsageFault(*)
SecureFault(*)
SVCall
DebugMonitor
PendSV
SysTick(**)
```



## Example: SysTick Handler in Rust

- Interrupts functions are identical to plain Rust functions
  - Return nothing and take no arguments

```
use cortex_m_rt::exception;

#[exception]
fn SysTick() {
    static mut COUNT: i32 = 0;

    // `COUNT` is safe to access and has type `&mut i32`
    *COUNT += 1;

    println!("{}", COUNT);
}
```



